

The background features a complex network of thin grey lines and dots, forming a web-like structure. Scattered throughout are various triangles of different sizes and orientations, some with solid black dots at their vertices. The overall aesthetic is technical and minimalist.


# **FFNN for computer vision**

Miguel Pérez  
David Burrell  
Ignacio Martínez

---

# 1st Case: Traffic Sign recognition

---



# Traffic Sign Recognition


This task involved developing a neural net that can classify images of traffic signs

Image preprocessing performed:

1. Images are **loaded** from a .csv then **rescaled, resized** and **stored** in a vector.
2. **Labels** are attached to their corresponding images
3. The data was **split** into two parts: **train** and **test**.
4. The **train** was divided again into two parts: **train** and **validation**.
5. The tensor values are **normalized** from the range 0 - 255 to 0 - 1
6. The **numerical values** representing each traffic sign are transformed into a binary categorical value.

Once this is done, the data is ready to be fed to the Neural Network

# Architecture



Input layer: 150528

1st Hidden Layer: 1000

2nd Hidden Layer: 750

Output Layer: 43

```
opt = optimizers.SGD(lr=0.0001, decay=1e-6, momentum=0.99, nesterov=True)

mlp.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])
mlp.summary()
```

```
data = mlp.fit(tr_signs,
               tr_labels,
               batch_size=25,
               epochs=150,
               verbose=1,
               validation_data=(va_signs, va_labels),
               callbacks=[tensorboard])
```


```
=====
Total params: 151,312,043
Trainable params: 151,312,043
Non-trainable params: 0
=====
```



# Results

ffNN took 0.19642043113708496 seconds

Test loss: 0.4035492646727205 - Accuracy: 0.9085872577828383





# **2nd Case:** **General Purpose Recognition**

---


# Designed NN

We developed a FFNN trying to achieve the best score (accuracy)

We performed some steps on several trials and selected the values that gave us better accuracy.

1. We modified **number of layers** and we changed the **size of each layer**.
2. We varied the **activation function** for the inner layers. The best one was the relu.
3. We tried adding **weight initializers**. However, none of them improved the accuracy.
4. We tested **batch normalization** to the output of each layer to normalize the output of each layer. Again, this measure did not improved the performance
5. We used **Dropout** regularization to reduce overfitting. We applied values of 0.5 and 0.2.
6. We used the following optimizer:
7. optimizers. **SGD**(lr=0.001, decay=1e-6, momentum=0.9. **nesterov=True**)
8. We used as loss function the **categorical crosentropy** and **accuracy** as the metric.

# Architecture



Input layer: 3072

1st Hidden Layer: 1000

2nd Hidden Layer: 750

Output Layer: 43

```
data = mlp.fit(x_train, tr_labels,  
              batch_size=32,  
              epochs=100,  
              verbose=1,  
              validation_data=(x_val, va_labels),  
              callbacks=[tensorboard])
```

```
opt = optimizers.SGD(lr=0.0001, decay=1e-6, momentum=0.99, nesterov=True)  
  
mlp.compile(optimizer=opt, loss='categorical_crossentropy', metrics=['accuracy'])  
mlp.summary()
```

```
=====
```

|                       |           |
|-----------------------|-----------|
| Total params:         | 3,898,850 |
| Trainable params:     | 3,898,850 |
| Non-trainable params: | 0         |


```
=====
```




# Results

```
ffNN took 0.5458638668060303 seconds  
Test loss: 3.007897315597534 - Accuracy: 0.2825
```

Took around 6 minutes and 30 seconds to train!





•

**Thanks for your attention!**